

carPredictionNotebook

July 29, 2025

1 Car Price Prediction Project

```
[29]: # Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report
```

```
[30]: # Define data
df = pd.read_csv(r"Assets\car_data_v2.csv")
print(df.head())
```

	car_name	brand	model	vehicle_age	km_driven	mileage \
0	Maruti Alto	Maruti	Alto	9	120000	19.70
1	Maruti Alto	Maruti	Alto	9	37000	20.92
2	Maruti Wagon R	Maruti	Wagon R	8	35000	18.90
3	Maruti Wagon R	Maruti	Wagon R	3	17512	20.51
4	Hyundai Venue	Hyundai	Venue	2	20000	18.15

	max_power	seats	selling_price
0	46.30	5	120000
1	67.10	5	226000
2	67.10	5	350000
3	67.04	5	410000
4	118.35	5	1050000

```
[31]: print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2119 entries, 0 to 2118
Data columns (total 9 columns):
```

```

#   Column      Non-Null Count  Dtype
---  -
0   car_name    2119 non-null    object
1   brand        2119 non-null    object
2   model        2119 non-null    object
3   vehicle_age  2119 non-null    int64
4   km_driven    2119 non-null    int64
5   mileage      2119 non-null    float64
6   max_power    2119 non-null    float64
7   seats        2119 non-null    int64
8   selling_price 2119 non-null    int64
dtypes: float64(2), int64(4), object(3)
memory usage: 149.1+ KB
None

```

```
[32]: print(df.describe())
```

```

      vehicle_age    km_driven    mileage    max_power    seats \
count  2119.000000    2119.000000  2119.000000  2119.000000  2119.000000
mean     6.153374   42207.621992    22.574856    61.802931    5.002832
std     3.524845   27950.561196     3.008683    13.112960    0.176503
min      0.000000     581.000000    14.400000    38.400000    4.000000
25%      4.000000   21000.000000    20.510000    53.260000    5.000000
50%      5.000000   38000.000000    22.740000    67.000000    5.000000
75%      8.000000   58494.000000    23.950000    67.050000    5.000000
max     29.000000  425785.000000    33.540000   123.370000    7.000000

      selling_price
count    2.119000e+03
mean     3.287744e+05
std      1.496699e+05
min      4.000000e+04
25%      2.490000e+05
50%      3.150000e+05
75%      3.900000e+05
max      1.240000e+06

```

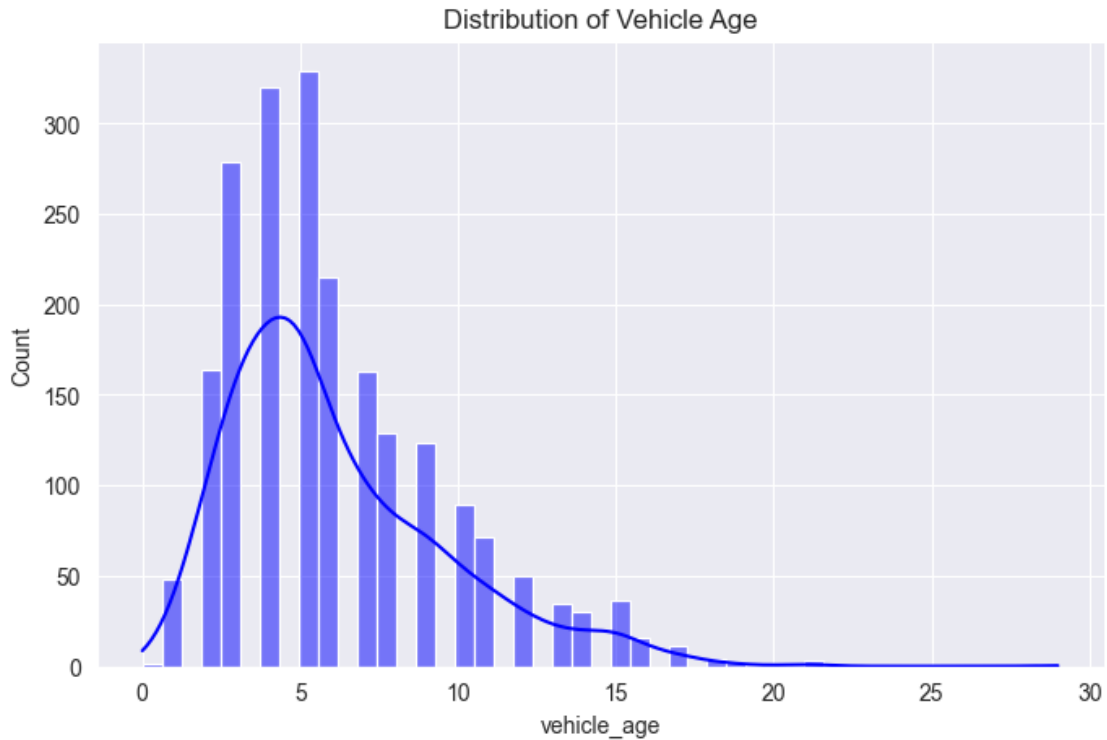
```
[33]: print(df.isnull().sum())
```

```

car_name      0
brand         0
model         0
vehicle_age   0
km_driven     0
mileage       0
max_power     0
seats         0
selling_price 0
dtype: int64

```


```
[34]: # Check the distribution of vehicle ages in order to identify trends related to
      ↪ vehicle age and pricing
plt.figure(figsize=(8,5))
sns.histplot(df['vehicle_age'], kde=True, color='blue')
plt.title('Distribution of Vehicle Age')
plt.show()
```




```
[35]: #Relationship between km_driven and selling_price
plt.figure(figsize=(8,5))
sns.scatterplot(x='km_driven', y='selling_price', data=df)
plt.title('Km Driven vs Selling Price')
plt.show()
```



```
[36]: # Relationship between power vs Selling Price
plt.figure(figsize=(8,5))
sns.scatterplot(x='max_power', y='selling_price', data=df)
plt.title('Max Power vs Selling Price')
plt.show()
```




```
[37]: # Brand Distribution of Selling Price
plt.figure(figsize=(12,8))
sns.boxplot(x='brand', y='selling_price', data=df)
plt.title('Brand vs Selling Price')
plt.xticks(rotation=45)
plt.show()
```

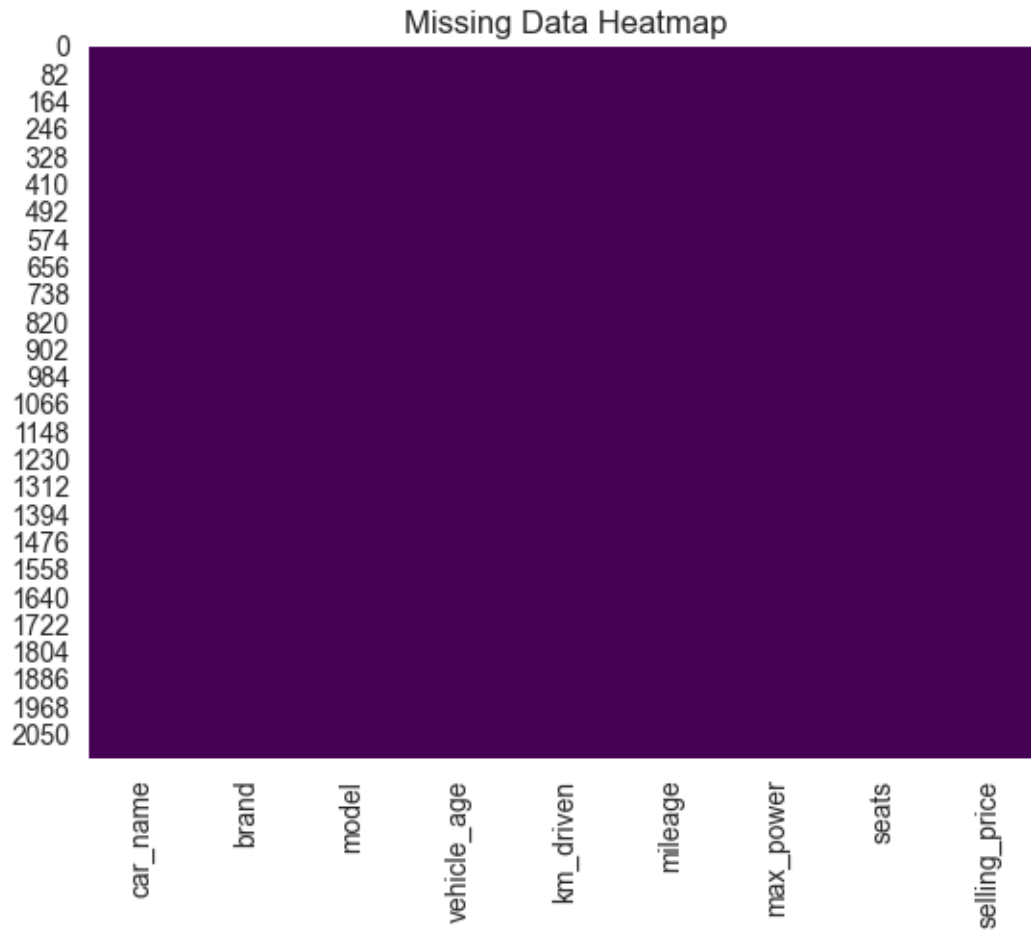


```
[38]: # Selection of the numeric columns for correlation analysis
numeric_columns = df.select_dtypes(include=['float64', 'int64'])

#Correlation matrix
plt.figure(figsize=(10,8))
sns.heatmap(numeric_columns.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix')
plt.show()
```



```
[39]: # Missing Value Check
sns.heatmap(df.isnull(), cbar=False, cmap='viridis')
plt.title('Missing Data Heatmap')
plt.show()
```



```
[40]: # Check for duplicate rows
df.duplicated().sum()

# # Fixing duplicate rows:
# df.drop_duplicates(inplace=True)
```

```
[40]: np.int64(34)
```

```
[41]: #data types
df.dtypes
```


```
[41]: car_name      object
brand      object
model      object
vehicle_age    int64
km_driven     int64
mileage      float64
max_power     float64
```



```
seats          int64
selling_price   int64
dtype: object
```

```
[42]: #Outliers
sns.boxplot(x=df['km_driven'], y=df['selling_price'], data=df)
```

```
[42]: <Axes: xlabel='km_driven', ylabel='selling_price'>
```



2 Modelling

LINEAR REGRESSION

```
[43]: # Feature selection
X = df.drop(['selling_price', 'car_name', 'model'], axis=1)
X = pd.get_dummies(X, drop_first=True)

y = df['selling_price']

# Train-test split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[44]: # Modeling the training data
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
```

```
[44]: LinearRegression()
```

```
[45]: # Evaluation
y_pred = lr_model.predict(X_test)
print("R2 Score:", r2_score(y_test, y_pred))
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))
```

R² Score: 0.8129121015140068
RMSE: 58325.83941323817

```
[46]: # Prediction of a new car's price
carInput = X_test.iloc[0:1]
predictedPrice = lr_model.predict(carInput)
print("Predicted Selling Price:", predictedPrice[0])
```

Predicted Selling Price: 399165.51263545104

LOGISTIC REGRESSION

```
[47]: medianPrice = df['selling_price'].median()
df['price_category'] = np.where(df['selling_price'] > medianPrice, 1, 0)
```

```
[48]: X = df.drop(['selling_price', 'price_category', 'car_name', 'model'], axis=1)
X = pd.get_dummies(X, drop_first=True)

y = df['price_category']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[49]: # Modeling of the training data
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_trainScaled = scaler.fit_transform(X_train)
X_testScaled = scaler.transform(X_test)

log_model = LogisticRegression(class_weight='balanced', max_iter=2000)
log_model.fit(X_trainScaled, y_train)
y_pred = log_model.predict(X_testScaled)
```

```
[50]: # Evaluation
y_pred = log_model.predict(X_test_scaled)

print("This is the Accuracy:", accuracy_score(y_test, y_pred))
print("This is the Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("The Classification Report:\n", classification_report(y_test, y_pred,
↳zero_division=0))
```

This is the Accuracy: 0.8608490566037735

This is the Confusion Matrix:

```
[[164  34]
 [ 25 201]]
```

The Classification Report:

	precision	recall	f1-score	support
0	0.87	0.83	0.85	198
1	0.86	0.89	0.87	226
accuracy			0.86	424
macro avg	0.86	0.86	0.86	424
weighted avg	0.86	0.86	0.86	424

```
[51]: # Scale the sample input (use the same scaler used for training)
carInput = X_test.iloc[0:1]
carInputScaled = scaler.transform(carInput)

# Predict if it will be expensive/affordable
predictedCategory = log_model.predict(carInputScaled)
print("Predicted Category (0=Affordable, 1=Expensive):", predictedCategory[0])
```

Predicted Category (0=Affordable, 1=Expensive): 1

```
[ ]:
```